

TMLE Under the Hood

Technical supplement for Segment G

Navigator → **nuisance fits** → **targeting** → **inference**

August 5, 2026

Andy Wilson · Kathryn Morrison

This Supplement Opens the Steps Hidden in the Live Demo

Live version

- One continuous Navigator-to-estimate story
- Compact `tmle()` call
- Core diagnostics and interpretation
- Script: `pneumonia_tmle_example.R`

Technical version

- Data generator and provenance check
- Super Learner fits for Q and g
- Clever covariate and targeting regression
- Efficient influence function and interval
- Script: `pneumonia_tmle_technical.R`

Both versions use the same frozen cohort, treatment, outcome, adjustment set, estimand, and workshop runtime settings.

Navigator Defines the Statistical Problem

Observed unit

$$O = (W, A, Y)$$

with $W = \{\text{age, prior pneumonia, prior vaccine}\}$.

Causal target

$$\Psi(P_0) = E_0[Y^{(1)} - Y^{(0)}].$$

Time order

1. Baseline history W
2. Baseline vaccination A
3. Pneumonia hospitalization Y by 12 months

The estimator is selected after this target and time order are fixed.

Identification Connects the Causal Target to Observed Data

Under consistency, conditional exchangeability, and positivity,

$$E\left[Y^{(a)}\right] = E_W\{E(Y \mid A = a, W)\}.$$


Therefore the identified risk-difference parameter is

$$\Psi(P_0) = E_W\{Q_0(1, W) - Q_0(0, W)\}, \quad Q_0(a, w) = E(Y \mid A = a, W = w).$$

Division of labor

The causal assumptions identify a statistical parameter. TMLE estimates that parameter from the observed data distribution.

The simcausal DAG Includes Baseline History and Treatment Selection



Age affects both prior-history variables; all three baseline variables affect vaccination and outcome.
Vaccination has a protective structural effect.

Generator Creates Confounding by Indication

Vaccination mechanism

$$\begin{aligned}\text{logit } P(A = 1 \mid W) = & -1.40 + 0.25 \frac{\text{age} - 65}{10} \\ & + 1.40 \text{priorPneumonia} \\ & + 1.20 \text{priorVaccine}.\end{aligned}$$

Higher-risk histories increase vaccination probability.

Outcome mechanism

$$\begin{aligned}\text{logit } P(Y = 1 \mid A, W) = & -3.00 - 0.80A \\ & + 0.25 \frac{\text{age} - 65}{10} \\ & + 1.60 \text{priorPneumonia} \\ & - 0.15 \text{priorVaccine}.\end{aligned}$$

The treatment coefficient is protective on the log-odds scale.

The observational treatment groups therefore differ before vaccination.

The Fixed Cohort Has a Reproducible Provenance Check

Canonical teaching data


- 5,000 observations
- 1,627 vaccinated
- 334 hospitalizations
- Frozen CSV verified row-for-row against the archived generator

Legacy RNG detail

- The original notebook set the seed before DAG plotting.
- Those setup steps advanced the RNG before `sim()`.
- The standalone generator reproduces that state and checks every row against the frozen CSV.

Cross-platform age values agree within 4.55×10^{-13} years; all discrete variables match exactly.

TMLE Uses Two Nuisance Fits and One Targeting Step



Q estimates conditional outcomes; g describes treatment assignment; the update solves the estimating equation associated with the selected risk-difference target.

Initial Outcome Regression Learns $Q(A, W)$

$$Q_0(a, w) = E(Y \mid A = a, W = w).$$

Fit the outcome model using observed A , W , and Y :

$$\hat{Q}(A_i, W_i).$$

These observed-treatment predictions are used in the targeting regression.

Predict each participant under both treatment strategies:

$$\hat{Q}(1, W_i), \quad \hat{Q}(0, W_i).$$

These counterfactual predictions are updated and averaged.

Predictions are bounded away from 0 and 1 before applying the logit update.

Super Learner Builds the Initial Nuisance Fits

Workshop library

- `SL.glm`: main-effects logistic regression
- `SL.glmnet`: penalized logistic regression
- `SL.ranger`: random forest

Ensembling

- Each candidate generates cross-validated predictions.
- The metalearner combines them to minimize prediction risk.
- The workshop uses three folds to keep the live run practical.

Flexible prediction can reduce nuisance-model misspecification, while the causal target and identification assumptions remain fixed upstream.

G-Computation Is the Initial Plug-In Contrast

Before targeting,

$$\hat{\Psi}_{\text{gcomp}} = \frac{1}{n} \sum_{i=1}^n \left\{ \hat{Q}(1, W_i) - \hat{Q}(0, W_i) \right\}.$$

In the transparent parametric comparator:

$$\hat{\Psi}_{\text{gcomp}} = -3.28 \text{ percentage points.}$$

This estimate relies directly on the outcome model. TMLE retains the standardization step but adds information from the treatment mechanism.

The Treatment Mechanism Learns $g(W)$

$$g_0(w) = P(A = 1 \mid W = w).$$


Role in diagnostics

- Describes treatment selection given measured history
- Reveals limited support and near-deterministic treatment
- Motivates redesign, restriction, or a revised target when needed

Role in estimation

- Defines inverse-probability weights
- Enters TMLE through the clever covariate
- Determines where outcome residuals receive greater leverage

Observed Data Show Shared Treatment Support



Parametric-GLM comparator

Range: 0.077 to 0.909

Outside $[0.05, 0.95]$: 0%

TMLE Super Learner g :
0.092 to 0.894.

Both are fitted-data support diagnostics; neither tests exchangeability or proves structural positivity.

Clever Covariate Focuses the Update

For a binary treatment and the average risk difference,

$$H(A, W) = \frac{A}{\hat{g}(W)} - \frac{1 - A}{1 - \hat{g}(W)}.$$

Under the two intervention values:

$$H(1, W) = \frac{1}{\hat{g}(W)}, \quad H(0, W) = -\frac{1}{1 - \hat{g}(W)}.$$

Observations receiving a treatment that was less probable given W contribute more strongly to the targeted residual correction.

A Logistic Fluctuation Estimates One Targeting Coefficient

Fit the intercept-free fluctuation model

$$\text{logit } E(Y \mid A, W) = \text{logit } \hat{Q}(A, W) + \epsilon H(A, W),$$

where $\text{logit } \hat{Q}(A, W)$ is an offset.

In the manual decomposition:

$$\hat{\epsilon} = 0.008503.$$

The update is low-dimensional and parameter-focused: the complex learning happened in the initial nuisance fits.

The Targeting Step Updates Both Intervention Predictions

$$\begin{aligned}\hat{Q}^*(1, W) &= \text{expit} \left[\text{logit } \hat{Q}(1, W) + \hat{\epsilon} \frac{1}{\hat{g}(W)} \right], \\ \hat{Q}^*(0, W) &= \text{expit} \left[\text{logit } \hat{Q}(0, W) - \hat{\epsilon} \frac{1}{1 - \hat{g}(W)} \right].\end{aligned}$$

Because the update is on the logit scale, all targeted predictions remain in the valid probability range.

The TMLE Is a Plug-In Contrast of Updated Risks

$$\hat{\Psi}_{\text{TMLE}} = \frac{1}{n} \sum_{i=1}^n \left\{ \hat{Q}^*(1, W_i) - \hat{Q}^*(0, W_i) \right\}.$$

Manual decomposition

−3.00 percentage points

tmle() implementation

−3.05 percentage points

Small differences reflect implementation details in nuisance fitting and targeting; both use the same data, adjustment set, and estimand.

Efficient Influence Function Supplies Uncertainty

For the average risk difference,

$$\begin{aligned} D^*(O) &= H(A, W)\{Y - Q^*(A, W)\} \\ &\quad + Q^*(1, W) - Q^*(0, W) - \Psi. \end{aligned}$$

Estimate

$$\widehat{\text{SE}}(\widehat{\Psi}) = \frac{\text{sd}\{\widehat{D}^*(O_i)\}}{\sqrt{n}}, \quad \widehat{\Psi} \pm 1.96 \widehat{\text{SE}}.$$

The manual implementation has $\overline{\widehat{D}^*} = -8.8 \times 10^{-13}$, a numerical check that the targeted estimating equation is solved. Its SE is 0.802 percentage points.

Manual Step 1: Fit Q and g

```
learners <- c("SL.glm", "SL.glmnet", "SL.ranger")

Q_fit <- SuperLearner(
  Y = Y, X = data.frame(A = A, W),
  family = binomial(), SL.library = learners,
  cvControl = list(V = 3L)
)
QAW <- Q_fit$SL.predict
Q1  <- predict(Q_fit, data.frame(A = 1, W))$pred
Q0  <- predict(Q_fit, data.frame(A = 0, W))$pred

g_fit <- SuperLearner(
  Y = A, X = W, family = binomial(),
  SL.library = learners, cvControl = list(V = 3L)
)
g <- pmin(pmax(g_fit$SL.predict, 0.025), 0.975)
H <- A / g - (1 - A) / (1 - g)
```

Manual Step 2: Target, Average, and Infer

```
fluctuation <- glm(  
  Y ~ -1 + H + offset(qlogis(QAW)),  
  family = binomial()  
)  
epsilon <- coef(fluctuation)[["H"]]  
  
Q1_star <- plogis(qlogis(Q1) + epsilon / g)  
Q0_star <- plogis(qlogis(Q0) - epsilon / (1 - g))  
QAW_star <- plogis(qlogis(QAW) + epsilon * H)  
  
psi <- mean(Q1_star - Q0_star)  
eif <- H * (Y - QAW_star) +  
  Q1_star - Q0_star - psi  
se <- sd(eif) / sqrt(length(Y))  
ci <- psi + qnorm(c(0.025, 0.975)) * se
```

The Package Implementation Encapsulates the Same Sequence

```
fit <- tmle(  
  Y = Y, A = A, W = W,  
  Q.SL.library = learners,  
  g.SL.library = learners,  
  V.Q = 3, V.g = 3  
)  
  
psi <- fit$estimates$ATE$psi  
ci <- fit$estimates$ATE$CI  
  
g1W <- fit$g$g1W  
Q0_initial <- fit$Qinit$Q[, "Q0W"]  
Q1_initial <- fit$Qinit$Q[, "Q1W"]  
Q0_updated <- fit$Qstar[, "Q0W"]  
Q1_updated <- fit$Qstar[, "Q1W"]
```


The manual version exposes the mechanics; the package version is the practical analysis route.

Technical Results Share One Risk-Difference Scale

Estimator	RD (pp)	95% lower	95% upper
Crude association	+0.76	–	–
Parametric g-computation	–3.28	–	–
IPTW	–2.94	–	–
Manual TMLE decomposition	–3.00	–4.57	–1.43
tmle() package	–3.05	–4.60	–1.50
DGP benchmark (frozen W)	–4.04	–	–

The benchmark averages model-implied counterfactual risks over the frozen covariate distribution; it would not be observed in an applied analysis.

Estimator Comparison Separates Association from Adjustment



The adjusted estimators recover a protective direction in this sample. Their similarity is useful for implementation review, while causal validity still depends on the shared study assumptions.

Double Robustness Uses Information from Both Nuisance Models

For a point-treatment average effect, TMLE is consistent under appropriate regularity and positivity conditions if either

$$\hat{Q} \longrightarrow Q_0 \quad \text{or} \quad \hat{g} \longrightarrow g_0.$$

Practical implication

- Fit both nuisance functions carefully.
- Inspect support and prediction behavior.
- Use flexible learners when justified.

Scope

- This is an estimator property.
- It does not identify unmeasured confounding.
- Finite-sample performance can still be poor.

Positivity Diagnostics Guide the Estimand and Analysis Population

When $\widehat{g}(W)$ approaches 0 or 1,

$$|H(A, W)|$$

can become large, increasing variance and sensitivity to nuisance-model error.

Analysis options

- Improve design or data collection
- Restrict to a population with shared support
- Use a treatment rule or overlap-weighted target

Truncation

- Can stabilize an estimator
- Introduces a bias-variance tradeoff
- Should be prespecified and sensitivity-analyzed

Influence-Curve Inference Requires Rate and Regularity Conditions

The usual Wald interval relies on an asymptotic linear expansion:

$$\hat{\Psi} - \Psi_0 = \frac{1}{n} \sum_{i=1}^n D^*(O_i) + o_p(n^{-1/2}).$$

A key remainder is controlled when the product of nuisance errors is small:

$$\|\hat{Q} - Q_0\| \|\hat{g} - g_0\| = o_p(n^{-1/2}).$$

Cross-fitting can reduce empirical-process restrictions for highly adaptive learners. The workshop script uses the standard single-fit tmle implementation with cross-validated Super Learner nuisance fits.

The Causal Interpretation Retains the Navigator Assumptions

Identification

- Exchangeability given measured W
- Positivity in the target population
- Consistency of treatment versions

Measurement

- Correct vaccination classification
- Correct hospitalization ascertainment

Time and selection

- Common eligibility and treatment time zero
- Complete or appropriately modeled follow-up
- Appropriate handling of missingness and selection

Transport

- The target population is the synthetic cohort
- No clinical generalization is intended

Audit Trail Connects Data, Code, Output, and Session State

Artifact	Role
data/pneumonia_data.csv	Frozen cohort used throughout the Navigator and Segment G
simulate_pneumonia_data.R	Structural equations, legacy-state reproduction, and value check
pneumonia_tmle_example.R	Short live analysis
pneumonia_tmle_technical.R	Manual decomposition, package fit, diagnostics, plots, and outputs
outputs/*.csv	Machine-readable estimates and diagnostics
outputs/session_info.txt	R and package environment

Rebuilding begins with the technical R script, then compiles each Beamer source twice.

TMLE, AIPW, and DML Use Orthogonal Scores

	TMLE	AIPW / one-step	DML
Core idea	Target an initial plug-in fit	Add an influence-function correction	Estimate an orthogonal score with sample splitting
Nuisance functions	Q and g	Q and g	Problem-specific nuisance functions
Parameter space	Plug-in estimate respects outcome bounds	One-step estimate need not be substitution-based	Depends on score and target
Inference	EIF-based under regularity/rate conditions	EIF-based under regularity/rate conditions	Orthogonal-score inference, typically cross-fitted

These are related frameworks, not interchangeable labels; implementation details and target parameters determine the exact guarantees.

Technical Takeaways

1. The Navigator fixes the population, time order, adjustment set, and risk-difference target before estimation.
2. TMLE begins with flexible estimates of $Q(A, W)$ and $g(W)$.
3. The clever covariate and logistic fluctuation target Q for the selected parameter.
4. Averaging $Q^*(1, W) - Q^*(0, W)$ gives a substitution estimate.
5. The efficient influence function supports the standard error and provides an estimating-equation diagnostic.
6. Double robustness and machine learning improve estimation under stated conditions; the causal interpretation retains the study assumptions.

Reproducible result

$\hat{\Psi}_{\text{TMLE}} = -3.05$ percentage points, 95% CI $[-4.60, -1.50]$, in synthetic teaching data.

Causal Inference References

- Hernan MA, Robins JM. *Causal Inference: What If*. Chapman & Hall/CRC; 2020.
- Petersen ML, Porter KE, Gruber S, Wang Y, van der Laan MJ. Diagnosing and responding to violations in the positivity assumption. *Statistical Methods in Medical Research*. 2012.
- Petersen ML, van der Laan MJ. Causal models and learning from data: integrating causal modeling and statistical estimation. *Epidemiology*. 2014.
- van der Laan MJ, Rose S. *Targeted Learning*. Springer; 2011.

Machine Learning References

- Chernozhukov V, et al. Double/debiased machine learning for treatment and structural parameters. *The Econometrics Journal*. 2018.
- Kennedy EH. Semiparametric doubly robust targeted double machine learning: a review. *arXiv:2203.06469*. 2023.
- van der Laan MJ, Polley EC, Hubbard AE. Super Learner. *Statistical Applications in Genetics and Molecular Biology*. 2007.